

CNN in practice

| 고려대학교 산업경영공학과
| 서덕성

INDEX


- Making the most of your data
 - Data augmentation
 - Transfer learning
- All about convolutions
 - How to arrange them
 - How to compute them fast
- Implementation details
 - GPU / CPU
 - Bottleneck
 - Floating procedure

Making the most of your data

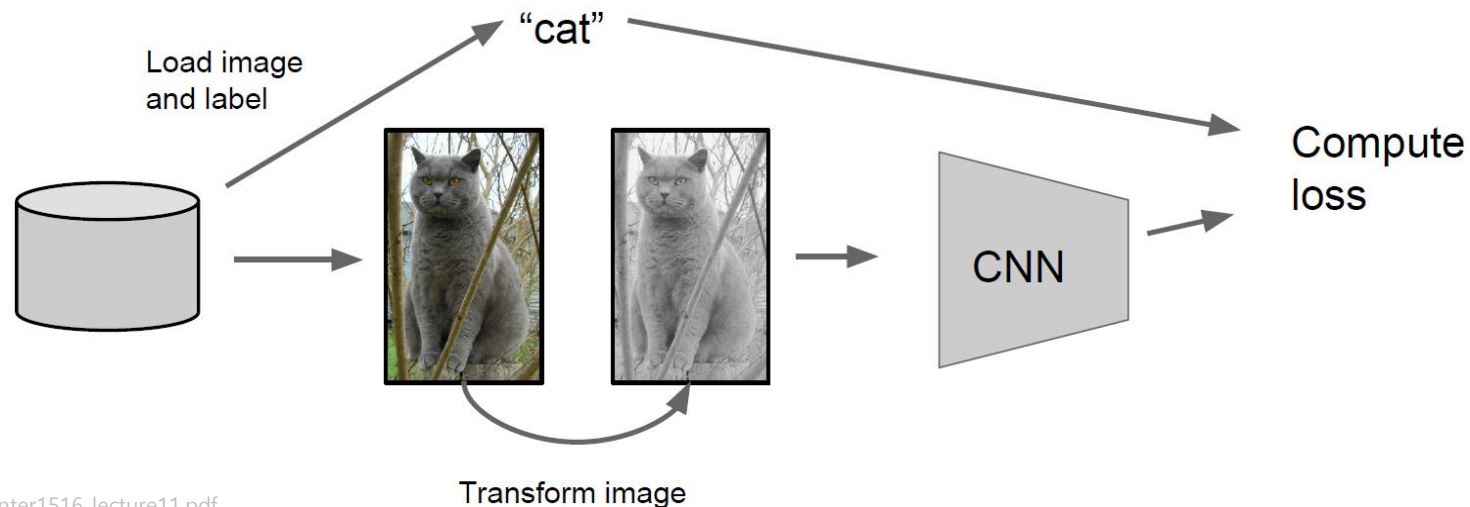
Data augmentation

❖ Original CNN

- 1개의 image와 1개의 label이 모델 학습시 input으로 들어감



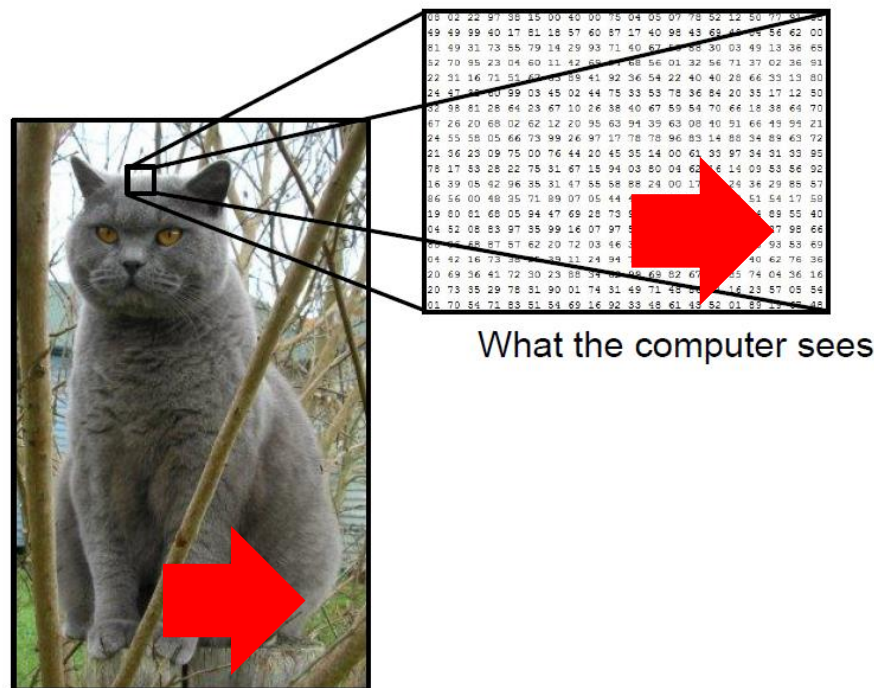
- 많은 image는 없지만, 많은 척 하고싶다 -> **Data augmentation**



Data augmentation

❖ Concept

- Label 변화 없이 픽셀값(독립변수)만 변화를 주자
- 변화 준 데이터까지도 학습에 사용하자
- 매우 많이 활용되고 있음



Data augmentation

01

❖ Horizontal flips

02

03

04

05

06

07





Data augmentation

❖ Random crops/scales

▪ ResNet 에서 사용한 방법

- Training : sample random crops / scales
 1. Pick random L in range [256, 480]
 2. Resize training image, short side = L
 3. Sample random 224 x 224 patch



Data augmentation

❖ Random crops/scales

▪ ResNet 에서 사용한 방법

• Training : sample random crops / scales

1. Pick random L in range $[256, 480]$
2. Resize training image, short side = L
3. Sample random 224×224 patch

• Test : average a fixed set of crops

1. Resize image at 5 scales : $\{224, 256, 384, 480, 640\}$
2. For each size, use 10 224×224 crops
: 4 corners + center, + flips


700



1000

224


224



Q




224



256

...



640

01

❖ Color jittering

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

02

- 가장 쉬운 방법 : 명도, 채도 등을 랜덤하게 변화시킴


03

04

05

06


07



- 상대적으로 복잡한 방법 : PCA를 이용
 1. [R, G, B] 3층을 가진 모든 training set의 pixel에 대해 PCA 적용
 2. PCA로부터 “color offset”을 추출
 3. 모든 픽셀에 대해 offset을 더함

❖ Color jittering


- 상대적으로 복잡한 방법 : PCA를 이용
 - [R, G, B] 3층을 가진 모든 training set의 pixel에 대해 PCA 적용
 - PCA로부터 “color offset”을 추출
 - 모든 픽셀에 대해 offset을 더해줌



$$\begin{bmatrix} I_{xy}^R \\ I_{xy}^G \\ I_{xy}^B \end{bmatrix} = \begin{bmatrix} I_{xy}^R \\ I_{xy}^G \\ I_{xy}^B \end{bmatrix} + [p_1 \ p_2 \ p_3] \begin{bmatrix} \alpha_1 \lambda_1 \\ \alpha_2 \lambda_2 \\ \alpha_3 \lambda_3 \end{bmatrix}$$

p_i : i-th 고유벡터
 λ_i : i-th 고유값
 α_i : i-th random variable (image마다 뽑음)
 $\alpha_i \sim N(0, 0.1^2)$


Color jittering




Data augmentation

❖ Get creative

- Random mix/combinations of
 - Translation
 - Rotation
 - Stretching
 - Shearing
 - Lens distortions ...(go crazy)




Data Augmentation



Dropout

❖ General theme

- Training : Add random noise (regularizer으로써 overfitting 막기 좋음)
 - Data augmentation
 - Dropout
 - DropConnect
 - Batch normalization
 - Model ensembles
- Test : Marginalize over the noise (종합)



DropConnect

Data augmentation

❖ Data augmentation summary

- Simple to implement, **use it**
- Especially useful for **small datasets**
- Fits into framework of **noise / marginalization**

Transfer learning

❖ CNN 학습시 한계

- 충분한 데이터가 없음
- 학습이 너무 오래 걸림 (VGGNet의 경우 2~3주)
- 최고 성능의 model을 internet에서 얻을 수 있음 -> 이를 활용하자!

❖ 상황별 Transfer learning 적용

- 데이터의 양
- 모델 학습시 사용한 데이터와 내가 가진 데이터의 유사성

	Similar	Difference
Very little data		
Quite a lot of data		

Transfer learning

❖ With VGGNet

VGGNet
(train on
imagenet)

image

conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

maxpool

FC-4096

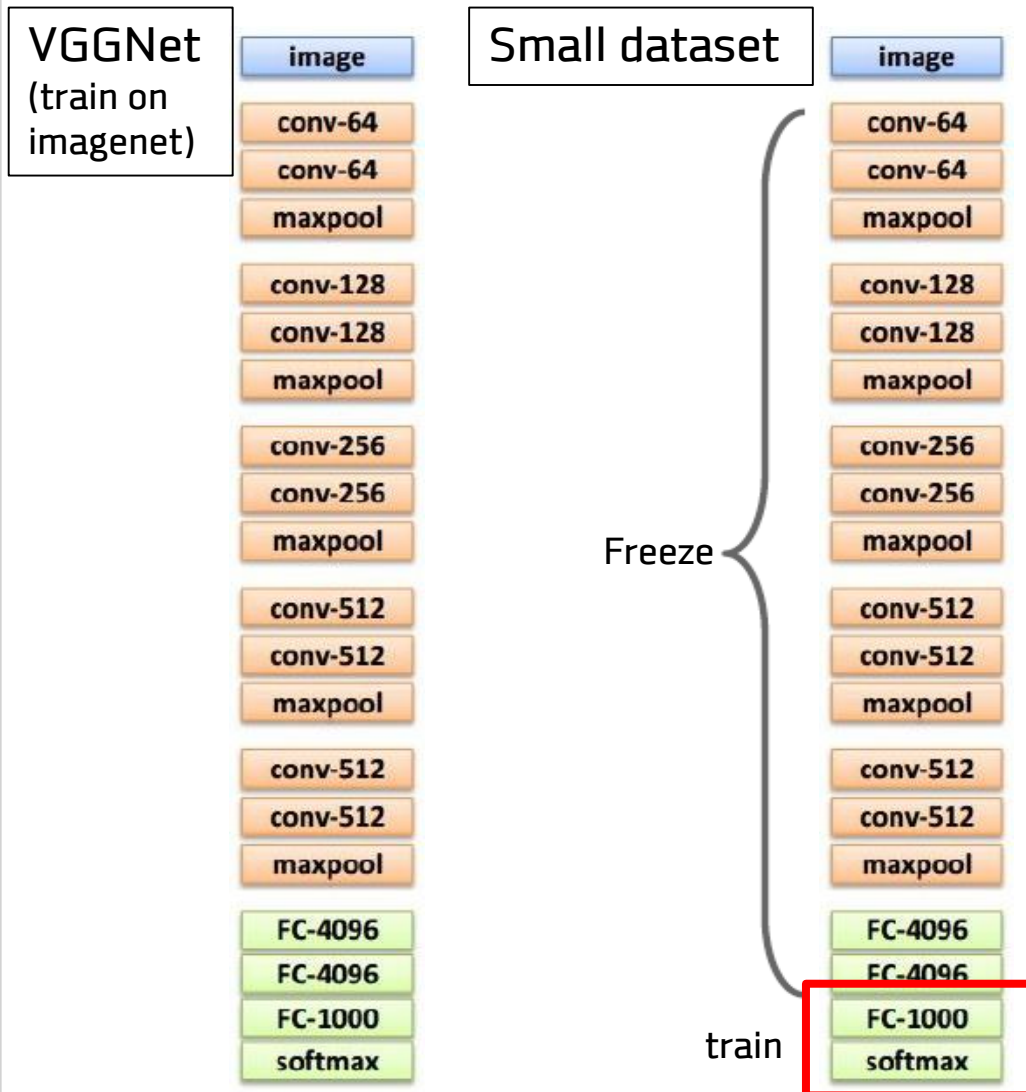
FC-4096

FC-1000

softmax

Transfer learning

❖ With VGGNet



Freeze 구간에서는 W , b 를 그대로 사용

- 즉, learning rate = 0
- 초반부 층에서는 점, 선, 색 등 이미지에 모두 이용가능한 feature가 있음
- 후반부 층에는 training image에 종속적인 feature가 있음
→ 유사한 데이터가 조금 있을 때 유용함

train 구간에서는 W , b 를 새로 만들

- 문제에 맞춰서 FC-10이나 SVM으로 변경


Transfer learning

❖ With VGGNet

VGGNet
(train on
imagenet)



Medium dataset



Train 구간 중 오렌지색(Conv 상위) 부분까지 학습

- 후반부 층에는 training image에 종속적인 feature가 있음
→ 나도 데이터가 좀 있으니 내 데이터에 적합하자

문제점 :

초반에는 큰 Loss로 인해 오렌지색 부분이 크게 흔들리는 경우가 있음

해결방법 :

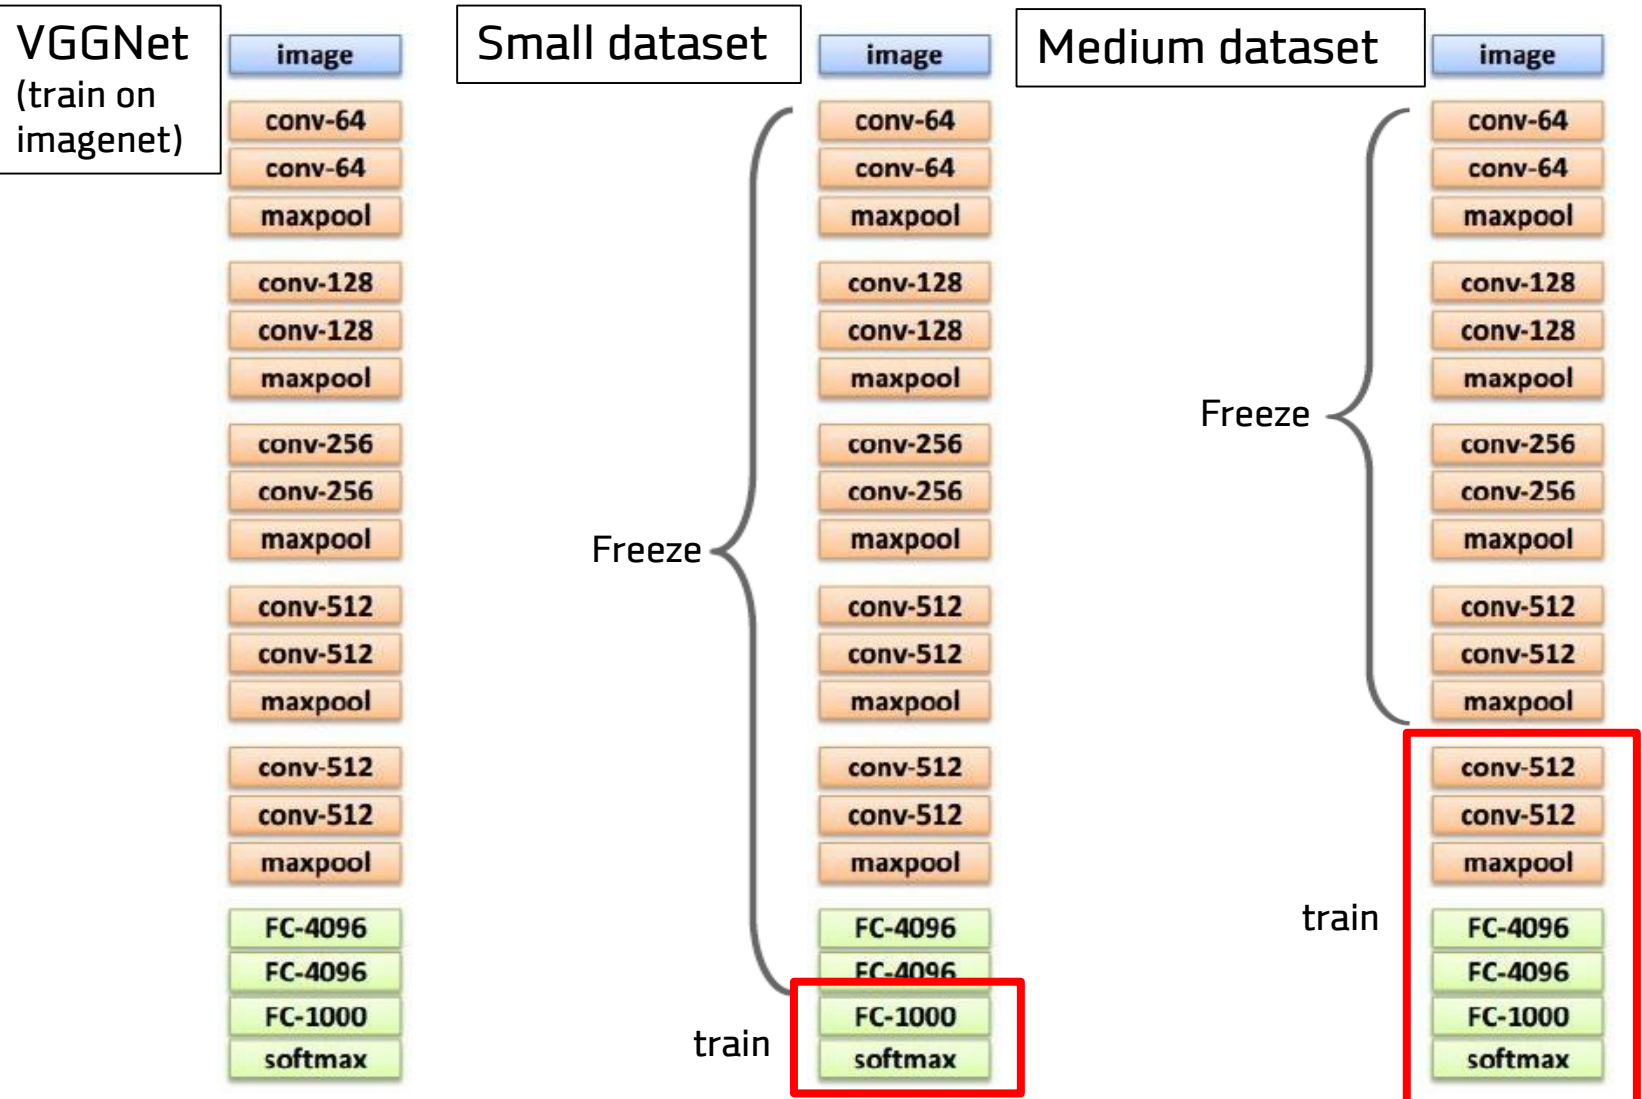
1. 오렌지 train부분을 잠시 freeze하고 초록색 부분만 학습한 뒤(Loss 좀 낮추고) 적절한 타이밍에 freeze를 풀어서 학습
2. (강의에서 준 tip)
원래 VGGNet의 learning rate값에 대해 초록색은 $\frac{1}{10}$, 오렌지색은 $\frac{1}{100}$ 을 주어 학습
→ 안정화 됨을 경험적으로 확인

Freeze

train

Transfer learning

❖ With VGGNet



Transfer learning

❖ 상황별 Transfer learning 적용

- 데이터의 양
- 모델 학습에 사용한 데이터와 내가 가진 데이터의 유사성

	Similar	Difference
Very little data	Use Linear classifier on top layer	In trouble...
Quite a lot of data	Finetune a few layers	Finetune a larger number of layers



Transfer learning

❖ 성능

- CNN Features off-the-shelf: an Astounding Baseline for Recognition

[Razavian et al, 2014]

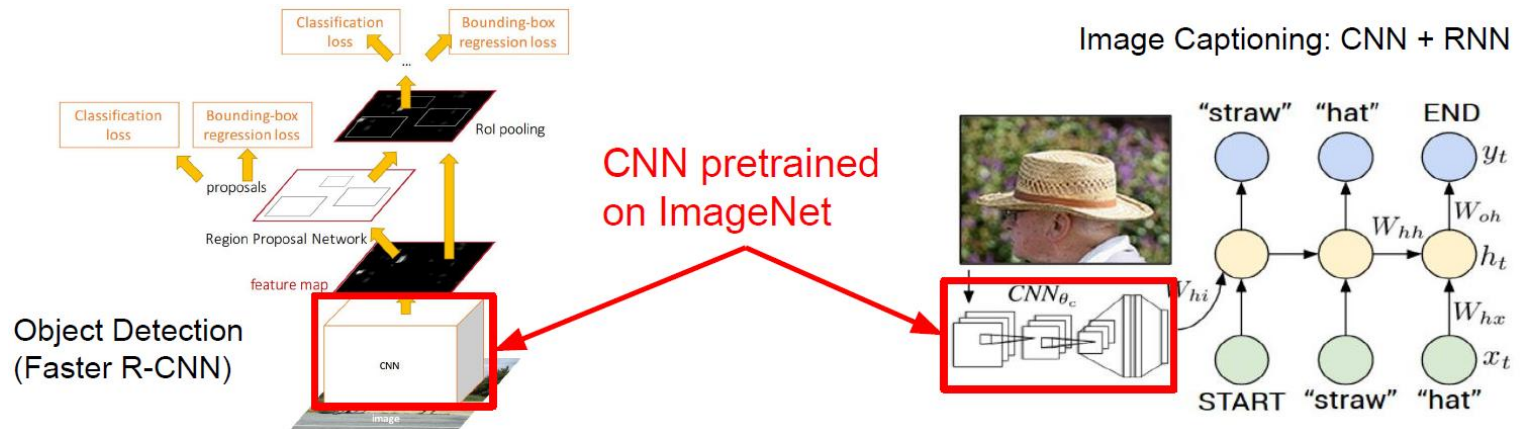
using simple augmentation techniques e.g. jittering. *The results strongly suggest that features obtained from deep learning with convolutional nets should be the primary candidate in most visual recognition tasks.*



Transfer learning

❖ 사용 예

- Object detection, Image captioning 등 CNN layer에서 활용 가능



Transfer learning

❖ Summary (have some dataset $< \sim 1\text{M}$)

- Find a very large dataset that has similar data, train a big ConvNet there
- Transfer learn to your dataset
- Caffe ConvNet library has a “Model Zoo” of pretrained models:

<https://github.com/BVLC/caffe/wiki/Model-Zoo>

Model Zoo

brandon-haystack edited this page 19 days ago · 102 revisions

Check out the [model zoo documentation](#) for details.

To acquire a model:

1. download the model gist by `./scripts/download_model_from_gist.sh <gist_id> <dirname>` to load the model metadata, architecture, solver configuration, and so on. (`<dirname>` is optional and defaults to `caffe/models`).
2. download the model weights by `./scripts/download_model_binary.py <model_dir>` where `<model_dir>` is the gist directory from the first step.

or visit the [model zoo documentation](#) for complete instructions.

Berkeley-trained models

- [Finetuning on Flickr Style](#): same as provided in `models/`, but listed here as a Gist for an example.
- BVLC GoogLeNet: `models/bvlc_googlenet`

Network in Network model

The Network in Network model is described in the following ICLR-2014 paper:

Network In Network
M. Lin, Q. Chen, S. Yan
International Conference on Learning Representations, 2014 (arXiv:1409.1556)

please cite the paper if you use the models.


All about convolutions

How to stack them

5 x 5 : 16

❖ Intro


- 3 x 3 filter를 2번 적용하면 1개 node는 몇 개의 pixel 정보를 포함할까?



How to stack them

❖ Intro

- 3 x 3 filter를 3번 적용하면 1개 node는 몇 개의 pixel 정보를 포함할까?




How to stack them

❖ The power of small filters


▪ Condition

- Input : $H \times W \times C$
- C 개의 filter를 사용해서 conv 지나도 다시 depth = C
- Stride = 1, padding 사용해서 H, W 유지




▪ One CONV with 7×7 filters

- Number of weights
 $= C \times (7 \times 7 \times C)$
 $= 49C^2$



▪ Three CONV with 3×3 filters

- Number of weights
 $= 3 \times [C \times (3 \times 3 \times C)]$
 $= 27C^2$




How to stack them

❖ The power of small filters


▪ Condition

- Input : $H \times W \times C$
- C 개의 filter를 사용해서 1층 지나도 다시 $\text{depth} = C$
- Stride = 1, padding 사용




▪ One CONV with 7x7 filters

- Number of multiply-adds
 $= (H \times W) \times C \times (7 \times 7 \times C)$
 $= 49HWC^2$



▪ Three CONV with 3x3 filters

- Number of multiply-adds
 $= (H \times W) \times 3 \times [C \times (3 \times 3 \times C)]$
 $= 27HWC^2$




How to stack them

❖ The power of small filters

▪ Condition

- Input : $H \times W \times C$
- C 개의 filter를 사용해서 1층 지나도 다시 $\text{depth} = C$
- Stride = 1, padding 사용



▪ One CONV with 7×7 filters

- Number of multiply-adds
 $= (H \times W) \times C \times (7 \times 7 \times C)$
 $= 49HWC^2$

▪ Three CONV with 3×3 filters

- Number of multiply-adds
 $= (H \times W) \times 3 \times [C \times (3 \times 3 \times C)]$
 $= 27HWC^2$

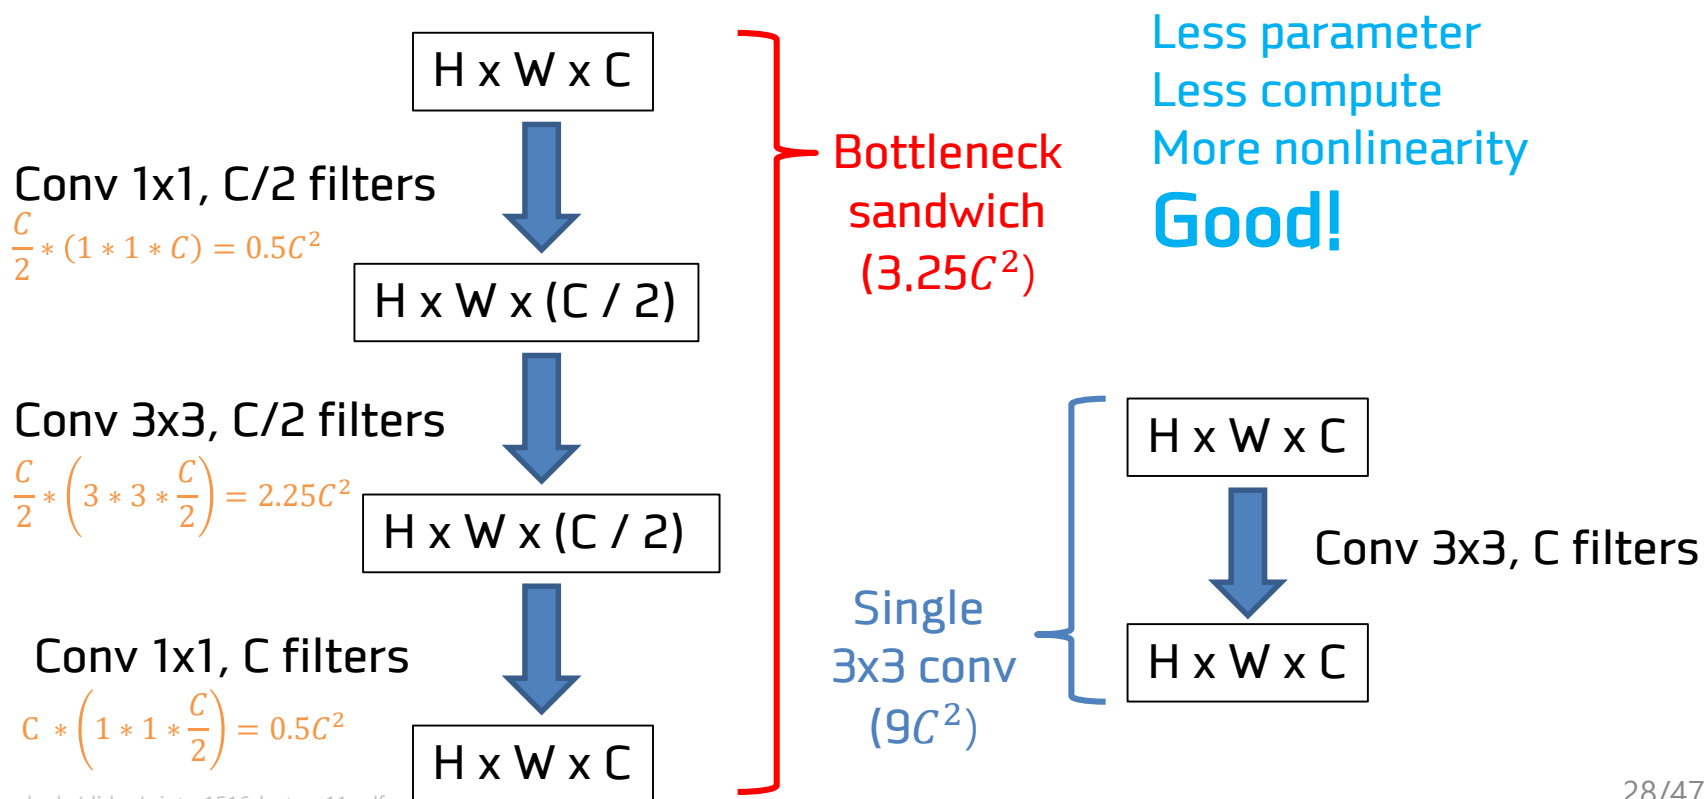
- Less compute, more nonlinearity = **GOOD!**

How to stack them

❖ Why not try 1 x 1 filters?

- Receptive 영역이 없어서 다른 pixel(input)까지 포괄하여 정보 저장 불가
 - 1 x 1 filter 단독으로는 절대 앞의 효과를 볼 수 없음

❖ 1 x 1을 활용하여 parameter 수를 줄이자 (Bottleneck sandwich)



How to stack them

❖ Still using 3 x 3 filters ... can we break it up?

- 정사각형 모양을 탈피하자
- 1 x 3 filter, 3 x 1 filter를 조합해서 3 x 3을 만듦

Conv 1x3, C filters

$$C * (1 * 3 * C) = 3C^2$$

H x W x C



H x W x C



H x W x C

Conv 3x1, C filters

$$C * (3 * 1 * C) = 3C^2$$

$$(6C^2)$$

Less parameter
Less compute
More nonlinearity

Good!

Single
3x3 conv
(9C²)

H x W x C




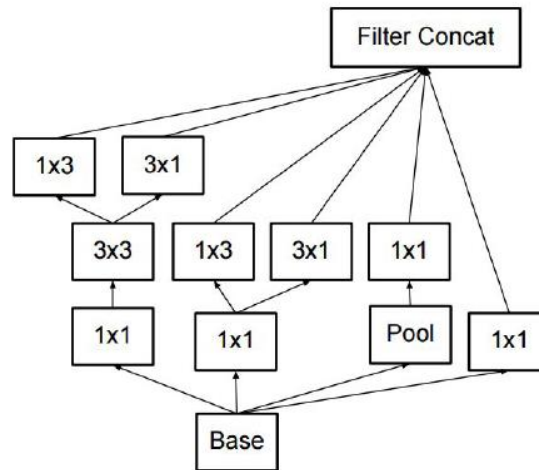

H x W x C

Conv 3x3, C filters

How to stack them

❖ Example

- Latest version of GoogLeNet incorporates all these ideas



Szegedy et al, "Rethinking the Inception Architecture for Computer Vision"

How to stack them

❖ Summary


- Replace large conv (5×5 , 7×7) with stacks of 3×3 conv
- 1×1 “bottleneck” conv are very efficient
- Can factor $N \times N$ conv into $1 \times N$ and $N \times 1$
- All of the above give
fewer parameters, less compute, more nonlinearity

How to compute them (fast)


❖ im2col

- 두 행렬의 원소단위 곱을 벡터단위 내적으로 접근 -> fast

Feature map: $H \times W \times C$





$K \times K \times C$ 크기의 벡터로 펼침
(벡터간 겹치는 부분이 많은 단점이 있음)
(경험적으로, 무시해도 되는 단점)



$K \times K \times C$ 크기의 벡터 N 개
 $(K^2C) * N$ matrix


Conv weights: D filters, each $K \times K \times C$




•

•

•



$K \times K \times C$ 크기의 벡터로 펼침




$K \times K \times C$ 크기의 벡터 D 개
 $D * (K^2C)$ matrix

행렬 곱



Output으로 나오는
feature map중 한 채널



$D \times N$ 의 행렬을 얻음
재조합 -> Feature map

How to compute them (fast)

❖ FFT (Fast Fourier Transform)

■ 이산 푸리에 변환(Discrete Fourier transform)

• 정의

– 이산적인 복소수 값 x_0, x_1, \dots, x_{N-1} 을 복소수 값 X_0, X_1, \dots, X_{N-1} 로 변환하는 식

$$- X_k = \sum_{n=0}^{N-1} x_n e^{-\frac{2\pi i}{N}kn}, \text{ (역연산)} x_k = \frac{1}{N} \sum_{n=0}^{N-1} X_n e^{\frac{2\pi i}{N}kn}, \quad k, n = 0, \dots, N-1$$

■ FFT

• 이산 푸리에 변환을 빠르게 수행하는 효율적인 알고리즘

• DFT는 $O(n^2)$ 의 연산이 필요. FFT는 $O(n \log n)$ 의 연산만으로도 가능

■ Convolution

• 정의

$$- \text{두 함수 } f, g \text{에 대해, } f * g(t) = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau = \int_{-\infty}^{\infty} f(t - \tau)g(\tau)d\tau$$

• Fourier transform 적용

$$- F(f * g) = F(f) * F(g)$$

■ 기존 연산 : $O(n^2)$ 의 연산. 본 접근법 : $O(n \log n)$ 의 연산

How to compute them (fast)

❖ FFT (Fast Fourier Transform)

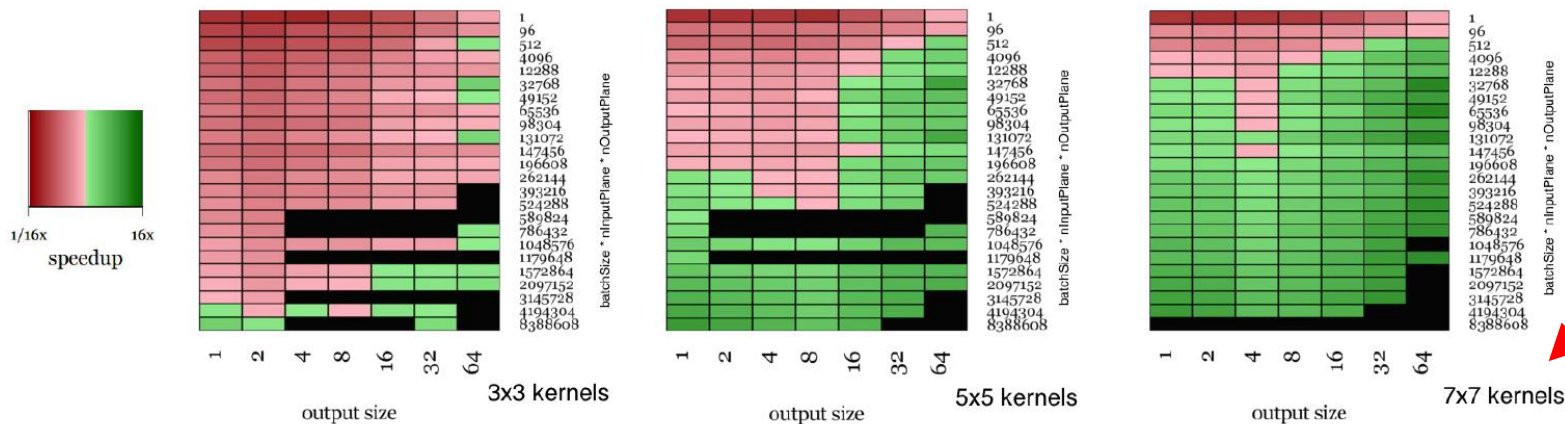
■ Implementing convolutions : FFT

- Compute FFT of weights : $F(W)$
- Compute FFT of image : $F(X)$
- Compute elementwise product : $F(W) * F(X)$
- Compute inverse FFT :

$$Y = F^{-1}(F(W) * F(X)) = F^{-1}(F(W * X)) = W * X$$

■ Performance

- Large filter에서만 좋은 성능



Vasilache et al, Fast Convolutional Nets With fbfft: A GPU Performance Evaluation

How to compute them (fast)

❖ Fast algorithms

- **Naive matrix multiplication** : Computing product of two $N \times N$ matrices takes $O(N^3)$ operations
- **Strassen's Algorithm** : Use clever arithmetic to reduce complexity to $O(N^{\log_2 7}) \sim O(N^{2.81})$


$$\begin{array}{lll} \mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{bmatrix} & \begin{array}{l} \mathbf{M}_1 := (\mathbf{A}_{1,1} + \mathbf{A}_{2,2})(\mathbf{B}_{1,1} + \mathbf{B}_{2,2}) \\ \mathbf{M}_2 := (\mathbf{A}_{2,1} + \mathbf{A}_{2,2})\mathbf{B}_{1,1} \\ \mathbf{M}_3 := \mathbf{A}_{1,1}(\mathbf{B}_{1,2} - \mathbf{B}_{2,2}) \\ \mathbf{M}_4 := \mathbf{A}_{2,2}(\mathbf{B}_{2,1} - \mathbf{B}_{1,1}) \\ \mathbf{M}_5 := (\mathbf{A}_{1,1} + \mathbf{A}_{1,2})\mathbf{B}_{2,2} \\ \mathbf{M}_6 := (\mathbf{A}_{2,1} - \mathbf{A}_{1,1})(\mathbf{B}_{1,1} + \mathbf{B}_{1,2}) \\ \mathbf{M}_7 := (\mathbf{A}_{1,2} - \mathbf{A}_{2,2})(\mathbf{B}_{2,1} + \mathbf{B}_{2,2}) \end{array} & \begin{array}{l} \mathbf{C}_{1,1} = \mathbf{M}_1 + \mathbf{M}_4 - \mathbf{M}_5 + \mathbf{M}_7 \\ \mathbf{C}_{1,2} = \mathbf{M}_3 + \mathbf{M}_5 \\ \mathbf{C}_{2,1} = \mathbf{M}_2 + \mathbf{M}_4 \\ \mathbf{C}_{2,2} = \mathbf{M}_1 - \mathbf{M}_2 + \mathbf{M}_3 + \mathbf{M}_6 \end{array} \\ \mathbf{B} = \begin{bmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{bmatrix} & & \\ \mathbf{C} = \begin{bmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{bmatrix} & & \end{array}$$

- 덧셈이 많아졌는데?
 - 큰 행렬의 경우, 행렬의 곱셈이 덧셈보다 더 많은 시간을 필요로 하기 때문에 덧셈을 더 하는 대신 곱셈을 덜 하는 것이 전체적으로 더 효율적이다.

How to compute them (fast)


❖ Fast algorithms

- 3 x 3 filter에 대해 fast algorithm을 적용
- 작은 filter에 대해 좋은 성능을 보임



N	cuDNN		F(2x2,3x3)		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	12.52	3.12	5.55	7.03	2.26X
2	20.36	3.83	9.89	7.89	2.06X
4	104.70	1.49	17.72	8.81	5.91X
8	241.21	1.29	33.11	9.43	7.28X
16	203.09	3.07	65.79	9.49	3.09X
32	237.05	5.27	132.36	9.43	1.79X
64	394.05	6.34	266.48	9.37	1.48X

Table 5. cuDNN versus $F(2 \times 2, 3 \times 3)$ performance on VGG Network E with fp32 data. Throughput is measured in Effective TFLOPS, the ratio of direct algorithm GFLOPS to run time.



N	cuDNN		F(2x2,3x3)		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	14.58	2.68	5.53	7.06	2.64X
2	20.94	3.73	9.83	7.94	2.13X
4	104.19	1.50	17.50	8.92	5.95X
8	241.87	1.29	32.61	9.57	7.42X
16	204.01	3.06	62.93	9.92	3.24X
32	236.13	5.29	123.12	10.14	1.92X
64	395.93	6.31	242.98	10.28	1.63X

Table 6. cuDNN versus $F(2 \times 2, 3 \times 3)$ performance on VGG Network E with fp16 data.

How to compute them (fast)

❖ Summary

- Im2col : Easy to implement, but big memory overhead
- FFT : Big speedups for small kernels
- “Fast Algorithms” seem promising, not widely used yet

Implementation details

GPU / CPU

❖ CPU / GPU


■ CPU

- Few, fast cores (1 - 16)
- Good at sequential processing

■ GPU

- Many, slower cores (thousands)
- Originally for graphics
- Good at parallel computation

■ GPUs are really good at matrix multiplication



← **GPU:** NVIDIA Tesla K40
with cuBLAS

← **CPU:** Intel E5-2697 v2
12 core @ 2.7 Ghz
with MKL



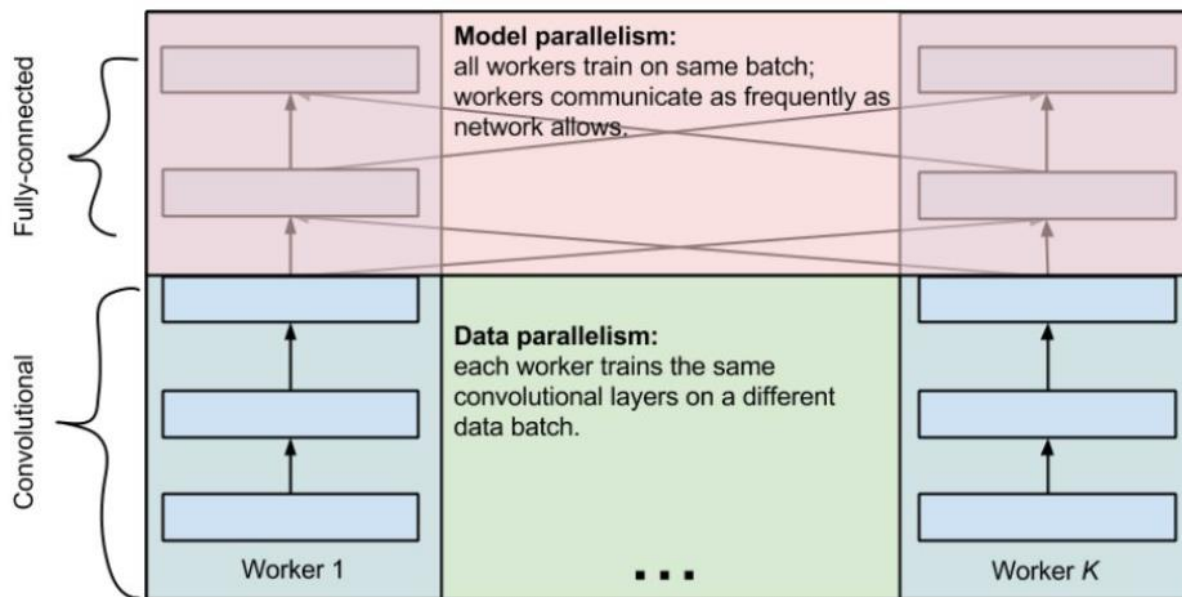
❖ GPU programming

- CUDA (NVIDIA only)
 - Write C code that runs directly on the GPU
 - Higher-level APIs: cuBLAS, cuFFT, **cuDNN**, etc
- OpenCL
 - Similar to CUDA, but runs on anything
 - Usually **slower** :(
- Udacity: Intro to Parallel Programming
 - <https://www.udacity.com/course/cs344>
 - For deep learning just use existing libraries

❖ 분산 학습

■ 개념


- 여러 개의 GPU(혹은 CPU)에 나눠서 학습시키자
- 다수의 GPU와 컴퓨터를 이용한 분산 학습을 지원한 프레임워크 출현
 - Google의 Tensorflow
 - Microsoft의 CNTK(Computational Network Toolkit)
- Multi-GPU training: More complex




GPU / CPU

❖ 분산 학습

- Google: Distributed CPU training




Data parallelism




Model parallelism

- Google: Synchronous vs Asynchronous



Synchronous Data Parallelism



Asynchronous Data Parallelism

Bottleneck

❖ 분산 학습

- GPU - CPU bottleneck
 - CPU data prefetch+augment thread running while GPU performs forward/backward pass
 - Allocate task efficiently
- CPU - disk bottleneck
 - SSD >> HDD
- GPU memory bottleneck
 - More GPU(expensive...)

Floating Point Precision

❖ 실수 표현 정밀도

- 컴퓨터는 bit 단위로 수를 표현함
 - CNN에서는 32bit single를 많이 사용함
- 추세
 - 32bit → 16bit
 - 16bit로 바뀌면서 실수 정밀도를 잃지만, 성능 손실을 최대한 줄이는 방법을 개발함 (Stochastic rounding, Gupta et al, "Deep Learning with Limited Numerical Precision", ICML 2015)
 - 이미 cuDNN에서 지원 중
 - Nervana fp16 kernels이 현재 가장 빠름

AlexNet (One Weird Trick paper) - Input 128x3x224x224



Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	92	29	62
CuDNN[R3]-fp16 (Torch)	cudnn.SpatialConvolution	96	30	66
CuDNN[R3]-fp32 (Torch)	cudnn.SpatialConvolution	96	32	64

OxfordNet [Model-A] - Input 64x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	529	167	362
Nervana-fp32	ConvLayer	590	180	410
CuDNN[R3]-fp16 (Torch)	cudnn.SpatialConvolution	615	179	436

GoogLeNet V1 - Input 128x3x224x224

Library	Class	Time (ms)	forward (ms)	backward (ms)
Nervana-fp16	ConvLayer	283	85	197
Nervana-fp32	ConvLayer	322	90	232
CuDNN[R3]-fp32 (Torch)	cudnn.SpatialConvolution	431	117	313




Floating Point Precision

❖ 실수 표현 정밀도

▪ Stochastic rounding

- 확률에 근거해 rounding
- Unbiased rounding scheme
- Expected rounding error is zero, $E(\text{Round}(x, \langle IL, FL \rangle)) = x$



- 정수부 x 자릿수
 - 예) $ab.c = abc.*2^{-1}$, $a.bc = abc.*2^{-2}$
- IL, FL : 정수부의 bit 수, 자릿수 담당 bit 수 (둘 다 설정해주는 값)
- ϵ : 매우 작은 양의 정수 2^{-FL}

$$\text{Round}(x, \langle IL, FL \rangle) = \begin{cases} \lfloor x \rfloor & \text{w.p. } 1 - \frac{x - \lfloor x \rfloor}{\epsilon} \\ \lfloor x \rfloor + \epsilon & \text{w.p. } \frac{x - \lfloor x \rfloor}{\epsilon} \end{cases}$$

Floating Point Precision

❖ 실수 표현 정밀도

▪ 도전

• 10bit → 1bit

- Courbariaux and Bengio, February 9 2016:
“BinaryNet: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1”
- 속도 향상을 위해 극한의 상황까지 도달
- Activation, weight = +1 or -1
- 빠른 계산
- Gradient를 구할 때는 조금 큰 bit를 사용

Floating Point Precision

❖ Summary

- GPUs much faster than CPUs
- Distributed training is sometimes used
 - Not needed for small problems
- Be aware of bottlenecks: CPU / GPU, CPU / disk
- Low precision makes things faster and still works
 - 32 bit is standard now, 16 bit soon
 - In the future: binary nets?

Q & A

01

02

03

04

05

06

07

